

# Retrieval-Augmented Generation with LlamaIndex



<https://github.com/danielbank/rag-llamaindex>

Daniel Bank

# A Brief Review

# What is Generative AI?

AI that can produce various types of content including text, images, audio, and data.

Examples:

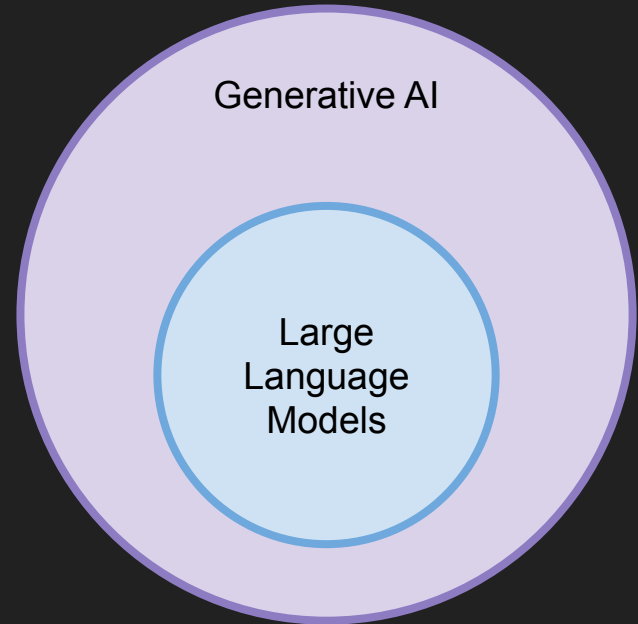
- RNNs (Recurrent Neural Networks / LSTM (Long Short-term Memory)
- GAN (Generative Adversarial Networks) [*generator and discriminator networks*]
- Transformer-based Models (BERT, GPT, etc)

# What is an LLM?

Large Language Models (LLM) are machine learning models that can comprehend and generate human language text.

Examples:

- GPT-4 (OpenAI) [*ChatGPT is the chat bot*]
- Gemini (Google)
- Claude 3 (Anthropic)



# What are the components of a prompt?

- **Instructions:** a specific task you want the model to perform.
- **Constraints:** restrictions on acceptable responses.
- **Input Data:** the data that we are interested to find a response for.
- **Output Indicator:** the type or format of the output.
- **Examples:** demonstrations that steer the model to better performance.
- **Context:** external information that shapes the response.

How do we teach an LLM  
about our data  
(data it wasn't trained on)  
?

# We could train the LLM on our data.

As seen in the [demo from the Getting Started with Generative AI APIs](#) presented at the February 2024 Phoenix AI Devs Meetup.

## Pros

- Consistent Output Quality
- Less Dependent on External Data

## Cons

- Expensive to Train
- Static Knowledge

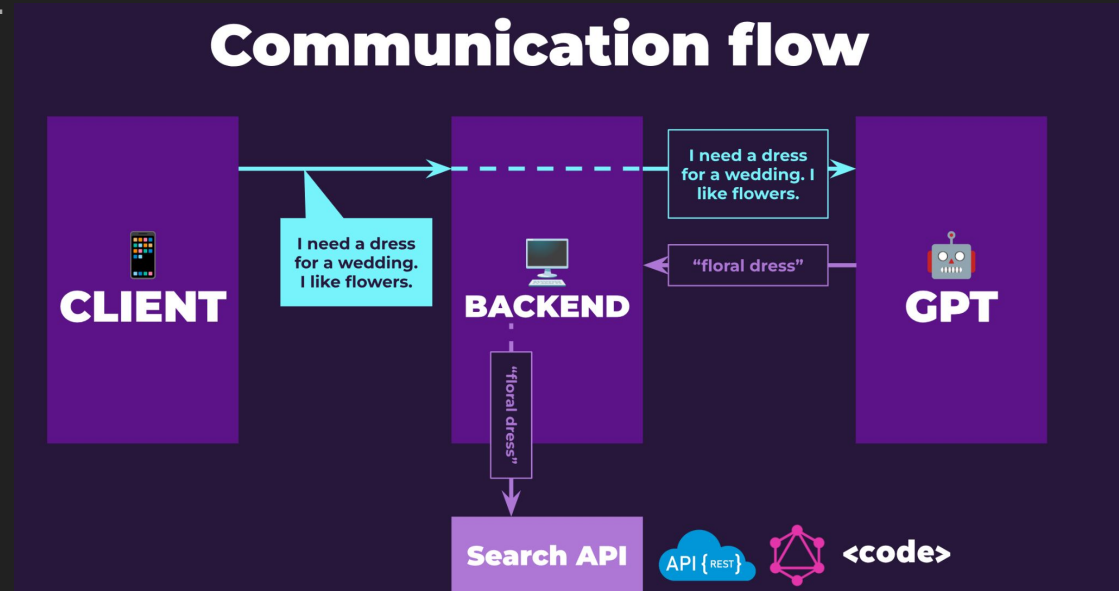
# We could leverage a tool outside the LLM.

As seen in [Ben Ilegbodu's GPT-Powered AI Shopping presentation](#) from HalfStack Phoenix. Here, the LLM translates unstructured input into structured input and then performs a server call with it.

NOTE:

LLM is not the “Frontend”

Backend acts as the intermediary





We could put include our data as context in the prompt (while still leveraging a tool outside the LLM).

- Original Paper: [Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks \(2021\)](#)
- Traditional RAG Pipeline Steps:
  1. Retrieve some context from the DB
  2. Stick that context into the LLM Prompt
  3. Call the LLM a single time to get a response
- DeepLearning.ai Course: [JavaScript RAG Web Apps with LlamaIndex](#)

# We could use Agentic RAG (outside the scope of this talk).

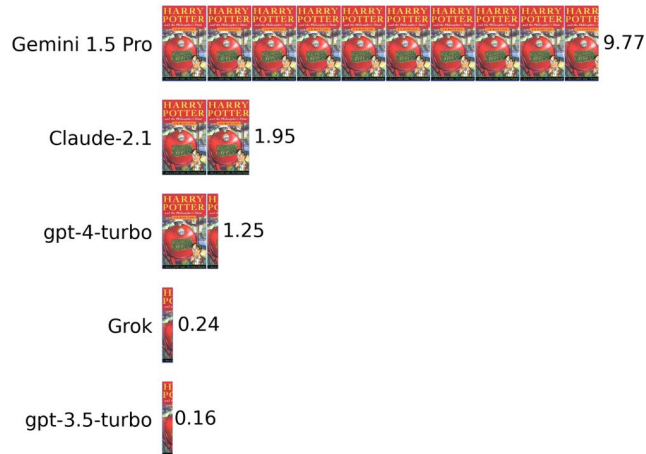
- Standard RAG Pipeline
  - Good for simpler questions over a small set of documents
- Agentic RAG Pipeline
  - Good for multiple step processing with possible dependencies of steps on each other
- How it works:
  - **Routing**: Add decision-making to route requests to multiple tools
  - **Tool Use**: Create an interface for agents to select a tool and provide the right arguments
  - **Multi-step reasoning with tool use**: Use an LLM to perform multi-step reasoning with a range of tools for retaining memory throughout that process
- DeepLearning.ai Course: [Building Agentic RAG with LlamaIndex](#)

# We could put our data as context in the prompt?

- Traditional RAG, you need to filter down the context you care about
- With a larger context window, you don't have to be as precise or prescriptive.

## Google's AI Gemini can fit almost 10 Harry Potter books in its context window

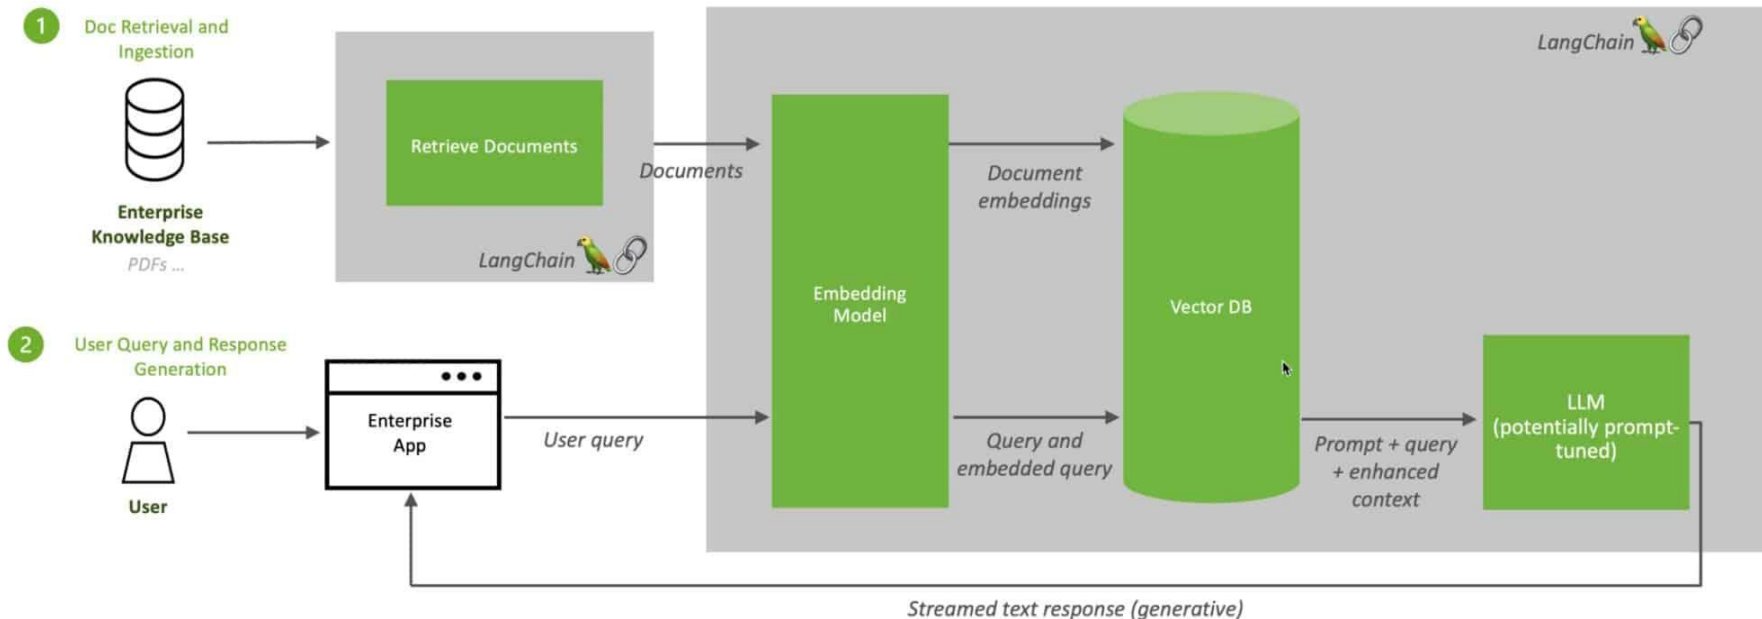
Number of copies of 'Harry Potter and the Sorcerer's Stone' books that can fit within each AI's context window



Let's Dive Deeper on RAG

# Traditional RAG Pipeline (from [NVIDIA's blog](#))

## Retrieval Augmented Generation (RAG) Sequence Diagram




# Retrieving Documents


- Get the data we intend to insert into the Vector DB in a well-defined shape
  - Example: We have text data in different file types (.md, .pdf, .csv, etc.)
- [LlamaParse](#) is a proprietary parsing service, exceptionally adept at transforming PDFs containing complex tables into a neatly organized markdown format:
  - What language to use for OCR?
  - Is there any text at different orientations?
  - Are there special instructions for the Embedding Model?

# Embedding Models

- An embedding model is a neural network that converts text into dense vector representations (embeddings).
- Different types of Embedding Models based on modality (text, image, audio)
- [Massive Text Embedding Benchmark \(MTEB\)](#)



# Vector Similarity (Dot Product)



$$\vec{A} \cdot \hat{B} = |A||B| \cos(\theta)$$

if the magnitude of B is 1, then...

$$C = \vec{A} \cdot \hat{B} = |A| \cos(\theta)$$



# Databases (A gentle introduction to Vector DBs)

- Relational Databases
  - When we think of databases, relational databases are usually what we picture
  - Great for Structured Data
  - Data is search by columns

ISBN	Year	Name	Author
0767908171	2003	A Short History of Nearly Everything	Bill Bryson
039516611X	1962	Silent Spring	Rachel Carson
0374332657	1998	Holes	Louis Sachar

# Vector Databases ([A gentle introduction to Vector DBs](#))

- Vast majority of data on the internet is unstructured (images, text, audio, video)
- Data is searched via content rather than keywords (similarity score, e.g. dot product)
  - Side note: PageRank (1998), the algorithm that powers Google Search, uses eigenvectors to determine page rank (similarity of pages to the query keywords)

Data UID <sup>1</sup>	Vector representation
00000000	[-0.31, 0.53, -0.18, ..., -0.16, -0.38]
00000001	[ 0.58, 0.25, 0.61, ..., -0.03, -0.31]
00000002	[-0.07, -0.53, -0.02, ..., -0.61, 0.59]

Let's build RAG with Milvus (a Vector DB) in  
Google Colab



# Milvus Bootcamp - Build RAG with Milvus Tutorial

- Milvus Bootcamp is a good resource to learn about the Milvus Vector DB:  
<https://github.com/milvus-io/bootcamp>
- Specifically, there is a Jupyter Notebook Tutorial for Building RAG with Milvus that we can run in a Google Colab:  
[https://github.com/milvus-io/bootcamp/blob/master/bootcamp/tutorials/quickstart/build\\_RAG\\_with\\_milvus.ipynb](https://github.com/milvus-io/bootcamp/blob/master/bootcamp/tutorials/quickstart/build_RAG_with_milvus.ipynb)

# Colab Tweaks

- Colab specifies a higher version of grpcio

```
! pip uninstall -y grpcio
```

```
! pip install grpcio==1.63.0
```

```
! pip install --upgrade pymilvus openai requests tqdm
```

- Don't put your **OPENAI\_API\_KEY** in the Notebook contents

```
import os
```

```
from google.colab import userdata
```


```
os.environ["OPENAI_API_KEY"] = userdata.get('OPENAI_API_KEY')
```



Let's build a  
RAG-Powered Web App  
with LlamaIndex and Next.js

# Demo Time

- GitHub Repo: <https://github.com/danielbank/rag-llamaindex>



# LlamaIndex - <https://github.com/run-llama/LlamaIndexTS>

- [Create LlamaIndex App](#) is a CLI Tool for Bootstrapping RAG Web Apps
- `npx create-llama@latest` - Run the CLI tool to bootstrap a web app
- `npm run generate` - Create the Document Embeddings for the knowledge-base data in `./data`
- `npm run dev` - Run the web app



# The Shape of the Data Matters

- Depending on the data, you may need to do pre-processing or use a beefier Embedding Model or more dimensions
- `BadRequestError: 400 This model's maximum context length is 8192 tokens, however you requested 77935 tokens (77935 in your prompt; 0 for the completion). Please reduce your prompt; or completion length.`

# The Kind of Data Matters (only text)

- The parser only supports text data in the quickstart example. If your ./data folder has any images, you will get an error
- **Error: Cannot calculate image nodes embedding without 'imageEmbedModel' set**

# Demo Time

- Let's replace the silly example with the codebase for [Histogramo](#), an Android App for detecting dice rolls

